

Haar Cascade, Face Detection

May 25, 2020

0.1 Face detection using OpenCV

One older (from around 2001), but still popular scheme for face detection is a Haar cascade classifier; these classifiers in the OpenCV library and use feature-based classification cascades that learn to isolate and detect faces in an image. You can read [the original paper proposing this approach here](#).

Let's see how face detection works on an example in this notebook.

```
In [1]: # import required libraries for this section
        %matplotlib inline
        import numpy as np
        import matplotlib.pyplot as plt
        import cv2

In [2]: # load in color image for face detection
        image = cv2.imread('images/multi_faces.jpg')

        # convert to RGB
        image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

        plt.figure(figsize=(20,10))
        plt.imshow(image)

Out[2]: <matplotlib.image.AxesImage at 0x7f48390f9160>
```



To use a face detector, we'll first convert the image from color to grayscale. For face detection this is perfectly fine to do as there is plenty non-color specific structure in the human face for our detector to learn on.

```
In [3]: # convert to grayscale
        gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

        plt.figure(figsize=(20,10))
        plt.imshow(gray, cmap='gray')
```

```
Out[3]: <matplotlib.image.AxesImage at 0x7f483909bef0>
```



Next we load in the fully trained architecture of the face detector, found in the file `detector_architectures/ haarcascade_frontalface_default.xml`, and use it on our image to find faces!

A note on parameters

How many faces are detected is determined by the function, `detectMultiScale` which aims to detect faces of varying sizes. The inputs to this function are: (image, scaleFactor, minNeighbors); you will often detect more faces with a smaller scaleFactor, and lower value for minNeighbors, but raising these values often produces better matches. Modify these values depending on your input image.

```
In [4]: # load in cascade classifier
        face_cascade = cv2.CascadeClassifier('detector_architectures/haarcascade_frontalface_def

        # run the detector on the grayscale image
        faces = face_cascade.detectMultiScale(gray, 4, 6)
```

The output of the classifier is an array of detections; coordinates that define the dimensions of a bounding box around each face. Note that this always outputs a bounding box that is square in dimension.

```
In [5]: # print out the detections found
print ('We found ' + str(len(faces)) + ' faces in this image')
print ("Their coordinates and lengths/widths are as follows")
print ('=====')
print (faces)
```

We found 13 faces in this image

Their coordinates and lengths/widths are as follows

=====

```
[[1295  94  96  96]
 [ 917 103  96  96]
 [1148 131  96  96]
 [ 683 149  96  96]
 [ 510 158  96  96]
 [1565 339  96  96]
 [ 588 390  96  96]
 [1157 391  96  96]
 [ 771 405  96  96]
 [ 383 414  96  96]
 [1345 411  96  96]
 [ 146 458  96  96]
 [ 996 526  96  96]]
```

Let's plot the corresponding detection boxes on our original image to see how well we've done.

```
In [ ]: img_with_detections = np.copy(image)    # make a copy of the original image to plot rectangles

# loop over our detections and draw their corresponding boxes on top of our original image
for (x,y,w,h) in faces:
    # draw next detection as a red rectangle on top of the original image.
    # Note: the fourth element (255,0,0) determines the color of the rectangle,
    # and the final argument (here set to 5) determines the width of the drawn rectangle
    cv2.rectangle(img_with_detections,(x,y),(x+w,y+h),(255,0,0),5)

# display the result
plt.figure(figsize=(20,10))
plt.imshow(img_with_detections)
```