

pool_visualization

May 26, 2020

0.1 # Pooling Layer

In this notebook, we add and visualize the output of a maxpooling layer in a CNN.

0.1.1 Import the image


```
In [1]: import cv2
import matplotlib.pyplot as plt
%matplotlib inline

# TODO: Feel free to try out your own images here by changing img_path
# to a file path to another image on your computer!
img_path = 'images/udacity_sdc.png'

# load color image
bgr_img = cv2.imread(img_path)
# convert to grayscale
gray_img = cv2.cvtColor(bgr_img, cv2.COLOR_BGR2GRAY)

# normalize, rescale entries to lie in [0,1]
gray_img = gray_img.astype("float32")/255

# plot image
plt.imshow(gray_img, cmap='gray')
plt.show()
```



0.1.2 Define and visualize the filters

In [2]: `import numpy as np`

```
## TODO: Feel free to modify the numbers here, to try out another filter!
filter_vals = np.array([[ -1,  -1,  1,  1], [ -1,  -1,  1,  1], [ -1,  -1,  1,  1], [ -1,  -1,  1,  1]])

print('Filter shape: ', filter_vals.shape)
```

Filter shape: (4, 4)

```
In [3]: # Defining four different filters,
        # all of which are linear combinations of the `filter_vals` defined above

        # define four filters
filter_1 = filter_vals
filter_2 = -filter_1
filter_3 = filter_1.T
filter_4 = -filter_3
filters = np.array([filter_1, filter_2, filter_3, filter_4])

        # For an example, print out the values of filter 1
print('Filter 1: \n', filter_1)
```

Filter 1:
[[-1 -1 1 1]

```
[-1 -1  1  1]
[-1 -1  1  1]
[-1 -1  1  1]]
```

0.1.3 Define convolutional and pooling layers

Initialize a convolutional layer so that it contains all your created filters. Then add a maxpooling layer, [documented here](#), with a kernel size of (4x4) so you can really see that the image resolution has been reduced after this step!

```
In [4]: import torch
import torch.nn as nn
import torch.nn.functional as F

# define a neural network with a convolutional layer with four filters
# AND a pooling layer of size (4, 4)
class Net(nn.Module):

    def __init__(self, weight):
        super(Net, self).__init__()
        # initializes the weights of the convolutional layer to be the weights of the 4
        k_height, k_width = weight.shape[2:]
        # assumes there are 4 grayscale filters
        self.conv = nn.Conv2d(1, 4, kernel_size=(k_height, k_width), bias=False)
        self.conv.weight = torch.nn.Parameter(weight)
        # define a pooling layer
        self.pool = nn.MaxPool2d(4, 4)

    def forward(self, x):
        # calculates the output of a convolutional layer
        # pre- and post-activation
        conv_x = self.conv(x)
        activated_x = F.relu(conv_x)

        # applies pooling layer
        pooled_x = self.pool(activated_x)

        # returns all layers
        return conv_x, activated_x, pooled_x

# instantiate the model and set the weights
weight = torch.from_numpy(filters).unsqueeze(1).type(torch.FloatTensor)
model = Net(weight)

# print out the layer in the network
print(model)
```

```

Net(
  (conv): Conv2d(1, 4, kernel_size=(4, 4), stride=(1, 1), bias=False)
  (pool): MaxPool2d(kernel_size=4, stride=4, padding=0, dilation=1, ceil_mode=False)
)

```

0.1.4 Visualize the output of each filter

First, we'll define a helper function, `viz_layer` that takes in a specific layer and number of filters (optional argument), and displays the output of that layer once an image has been passed through.

```

In [5]: # helper function for visualizing the output of a given layer
        # default number of filters is 4
        def viz_layer(layer, n_filters= 4):
            fig = plt.figure(figsize=(20, 20))

            for i in range(n_filters):
                ax = fig.add_subplot(1, n_filters, i+1, xticks=[], yticks=[])
                # grab layer outputs
                ax.imshow(np.squeeze(layer[0,i].data.numpy()), cmap='gray')
                ax.set_title('Output %s' % str(i+1))

```

Let's look at the output of a convolutional layer after a ReLu activation function is applied.

```

In [6]: # plot original image
        plt.imshow(gray_img, cmap='gray')





        # visualize all filters
        fig = plt.figure(figsize=(12, 6))
        fig.subplots_adjust(left=0, right=1.5, bottom=0.8, top=1, hspace=0.05, wspace=0.05)
        for i in range(4):
            ax = fig.add_subplot(1, 4, i+1, xticks=[], yticks=[])
            ax.imshow(filters[i], cmap='gray')
            ax.set_title('Filter %s' % str(i+1))

        # convert the image into an input Tensor
        gray_img_tensor = torch.from_numpy(gray_img).unsqueeze(0).unsqueeze(1)

        # get all the layers
        conv_layer, activated_layer, pooled_layer = model(gray_img_tensor)

        # visualize the output of the activated conv layer
        viz_layer(activated_layer)

```



0.1.5 Visualize the output of the pooling layer

Then, take a look at the output of a pooling layer. The pooling layer takes as input the feature maps pictured above and reduces the dimensionality of those maps, by some pooling factor, by constructing a new, smaller image of only the maximum (brightest) values in a given kernel area.

```
In [7]: # visualize the output of the pooling layer
        viz_layer(pooled_layer)
```

Output 1



Output 2



Output 3



Output 4



In []: